



Transactions and Isolation Levels

.NET

Transactions specify an **isolation level** that defines the degree to which one transaction must be isolated from a resource or from data modifications made by other transactions. **Isolation levels** are described in terms of which concurrency side effects are allowed.

[HTTPS://DOCS.MICROSOFT.COM/EN-US/SQL/CONNECT/JDBC/UNDERSTANDING-ISOLATION-LEVELS?VIEW=SQL-SERVER-VER15](https://docs.microsoft.com/en-us/sql/connect/jdbc/understanding-isolation-levels?view=sql-server-ver15)

Isolation Level and Read Errors

<https://www.geeksforgeeks.org/transaction-isolation-levels-dbms/#:~:text=>

Isolation levels define the degree to which a transaction must be isolated from other data modifications made by any other transaction. A **transaction isolation level** is determined by its permissiveness of the following.

- **Dirty Read** – When a transaction reads data that is new and has not yet been committed by another transaction.
- **Non-Repeatable read** – When a transaction reads the same row twice and gets a different value each time.
- **Phantom Read** – When two identical queries are executed but the rows retrieved by the two are different.

Isolation Levels

<https://docs.microsoft.com/en-us/sql/connect/jdbc/understanding-isolation-levels?view=sql-server-ver15>

Transaction isolation levels control the following:

- If **locks**, and what type, are enacted when data is read.
- How long read-locks are held.
- If a read operation references rows modified by another transaction:
 - Block until the exclusive lock on the row is freed.
 - Retrieve the committed version of the row that existed at the time the statement or transaction started.
 - Read the uncommitted data modification.


Choosing a transaction isolation level doesn't affect the locks that are acquired to protect data modifications. A transaction always gets an exclusive lock on any data it modifies and holds that lock until the transaction completes, regardless of its isolation level.

Isolation Levels – Serializable

<https://sqlperformance.com/2015/04/t-sql-queries/the-read-uncommitted-isolation-level>

The only transaction isolation level that provides complete isolation from concurrency effects is ***serializable***. The SQL Server implementation of the serializable isolation level means a transaction will see the latest committed data. The set of data encountered under ***serializable*** isolation is guaranteed not to change before the transaction ends.

| Isolation Level | Dirty Read | Non Repeatable Read | Phantom |
|-----------------|------------|---------------------|---------|
| Serializable | No | No | No |




Isolation Levels – Repeatable-Read

<https://sqlperformance.com/2015/04/t-sql-queries/the-read-uncommitted-isolation-level>

The *Repeatable Read* isolation level is guaranteed to read only committed data. If two identical queries are enacted simultaneously, they may return different results.

| Isolation Level | Dirty Read | Non Repeatable Read | Phantom |
|-----------------|------------|---------------------|---------|
| Repeatable read | No | No | Yes |




Isolation Levels – Read-Committed

<https://sqlperformance.com/2015/04/t-sql-queries/the-read-uncommitted-isolation-level>

Read committed can see committed data from different points in time – even for a single row if, for example, the query plan uses techniques like index intersection.

| Isolation Level | Dirty Read | Non Repeatable Read | Phantom |
|-----------------|------------|---------------------|---------|
| Read committed | No | Yes | Yes |




Isolation Levels – Read-Uncommitted

<https://sqlperformance.com/2015/04/t-sql-queries/the-read-uncommitted-isolation-level>

Read uncommitted is the weakest of the four transaction isolation levels. It allows all three "concurrency phenomena"; *dirty reads*, *non-repeatable reads*, and *phantom reads*.

| Isolation Level | Dirty Read | Non Repeatable Read | Phantom |
|------------------|------------|---------------------|---------|
| Read uncommitted | Yes | Yes | Yes |



Isolation Levels – In context

<https://docs.microsoft.com/en-us/sql/connect/jdbc/understanding-isolation-levels?view=sql-server-ver15#remarks>

The following table shows the concurrency side effects allowed by the different isolation levels.

| Isolation Level | Dirty Read | Non Repeatable Read | Phantom |
|------------------|------------|---------------------|---------|
| Read uncommitted | Yes | Yes | Yes |
| Read committed | No | Yes | Yes |
| Repeatable read | No | No | Yes |
| Snapshot | No | No | No |
| Serializable | No | No | No |

Transactions – Commit/Rollback/Savepoint

<https://www.techopedia.com/definition/16/commit#:~:text=>

- **Commit** - Refers to the saving of data permanently after a set of tentative changes. A commit ends a transaction within a Relational Database and allows all other users to see the changes.
- **Rollback** - In database technologies, a rollback is an operation which returns the database to some previous state. Rollbacks are important for database integrity, because they mean that the database can be restored to a clean copy even after erroneous operations are performed.
- **Savepoint** - A way of implementing nested transactions within a Relational Database Management System by indicating a point within a transaction that can be "rolled back to" without affecting any work done in the transaction before the **Savepoint** was created. Multiple **Savepoints** can exist within a single transaction.